DataQuest 2026

End-to-End Credit Risk Scorecard Framework

Champion vs. Baseline vs. Cost-Sensitive Modelling, Feature Engineering & Business Optimisation

Ditiro Letsoalo
Master of Science Candidate in Statistics and Data Science
University of Cape Town
View Reproducible Code on GitHub

May 22, 2026

Abstract

This report presents a complete, regulatory-grade credit risk framework developed for the DataQuest 2026 hackathon. Using a simulated retail loan portfolio of 120,960 applications, we address intensive structural missingness through grouped conditional imputation, normalise 14 formatting variants in categorical fields, and handle three inconsistent date-format profiles present in a single column. A full Weight of Evidence (WoE) and Information Value (IV) engine identifies the strongest predictors and confirms that several features carry no meaningful signal and should be excluded. The Champion logistic regression model, enhanced with targeted feature engineering, natural cubic splines, and interaction terms, achieves an AUC of **0.7977** and a Gini coefficient of **0.5954** on the held-out test set — a substantial lift over the provided baseline of 0.68. We additionally evaluate a **Cost-Sensitive Weighted** logistic regression that upweights minority-class (default) observations during training to address class imbalance. After Youden's J threshold calibration, both models converge to virtually identical discrimination performance, and the unweighted Champion is selected as the production model on grounds of regulatory simplicity and superior raw-count performance on actual defaults. Raw predicted probabilities are scaled into a production-ready point-based scorecard (base 600, PDO 20), and a business decision dashboard demonstrates how the approval threshold should be tuned to balance credit-loss risk against origination volume. The Population Stability Index (PSI) confirms score stability between training and test populations.

Contents

1 Executive Summary and Modelling Strategy

The core challenge of this project was to push logistic regression — an inherently linear, transparent model — as close as possible to the discriminatory ceiling set by a reference LightGBM (AUC 0.82) while retaining full interpretability for risk managers, business stakeholders, and regulators. This is achieved not through model complexity, but through *feature engineering quality* and *domain-informed data preparation.*

The workflow is structured in four sequential stages:

1. **Data Quality and Cleaning.** Resolve structural messiness: inconsistent capitalisation and synonyms in categorical variables, three different date formats in one column, and three distinct missingness mechanisms.
2. **Exploratory Data Analysis (EDA).** Use WoE / IV analysis, distribution visualisations, and bivariate interaction mapping to discover which features carry genuine predictive signal and in what form.
3. **Feature Engineering and Modelling.** Translate every EDA insight into a concrete model input: log transforms, interaction terms, binary flag features, and non-linear splines. Three competing models are evaluated: a raw-feature Baseline, a fully-engineered Champion, and a Cost-Sensitive Weighted variant.
4. **Business Deployment.** Convert the selected model into an actionable scorecard, calculate the PSI for stability monitoring, and build a threshold-optimisation tool that shows the volume–risk trade-off.

Result

Champion Test AUC: **0.7977** Gini: **0.5954** KS: **0.412** vs. published baseline: **0.68**. Improvement: **+11.8 percentage points** in absolute AUC.

Cost-Sensitive Weighted AUC: **0.7981** (Youden threshold: 0.51) — functionally identical to Champion; Champion selected as production model.

2 Research: Theoretical Foundations

2.1 Generalised Linear Models vs. Non-Linear Models for Classification

A Generalised Linear Model (GLM) extends ordinary regression to allow the response to follow distributions beyond the normal family. For binary classification, the natural choice is the **Binomial GLM** — logistic regression — which links the probability of the event to a linear predictor through the logit (log-odds) function:

$$\text{logit}(p_i) = \ln\left(\frac{p_i}{1-p_i}\right) = \eta_i = \beta_0 + \beta_1 x_{i1} + \cdots + \beta_k x_{ik} \tag{1}$$

The log-odds model is directly invertible to a probability:

$$p_i = \frac{e^{\eta_i}}{1+e^{\eta_i}} = \sigma(\eta_i) \tag{2}$$

Non-linear models such as Random Forests, Gradient Boosted Trees (LightGBM/XGBoost), and Neural Networks operate differently: they learn arbitrary, non-parametric mappings from features to outcomes by partitioning feature space into complex regions (trees) or through hierarchical latent representations (neural nets). Table 1 summarises the key trade-offs.

Table 1: GLM vs. Non-Linear Model Comparison for Credit Scoring

Property	**Logistic Regression (GLM)**	**Non-Linear (e.g. LightGBM)**
Interpretability	Each coefficient has a direct log-odds meaning	Requires SHAP or surrogate models to explain
Regulatory acceptance	High — SR 11-7 / Basel III model risk compliance	Requires additional model validation steps
Feature engineering requirement	High — must encode non-linearities manually	Low — trees learn splits automatically
Risk of overfitting	Low on tabular data	Moderate — requires careful tuning
Scorecard conversion	Direct log-odds scaling (PDO system)	Not directly applicable
Typical AUC on credit data	0.68–0.78 with good engineering	0.78–0.85

Key Insight

The business constraint — *logistic regression only* — is not a limitation but a feature. A model that a regulator and a credit manager can interrogate coefficient-by-coefficient is far more deployable than a black box, even if the black box achieves a slightly higher AUC.

2.2 The Interpretability–Complexity Trade-off

This trade-off is fundamental to applied machine learning. High-complexity models (gradient boosting, neural networks) generally achieve higher predictive performance but at the cost of transparency: it becomes impossible to tell *why* a particular customer received a low score. This matters enormously in credit:

- Regulators in South Africa (National Credit Act) and internationally (Basel III SR 11-7) require that an institution can *explain* any adverse credit decision to the applicant and to auditors.
- Model risk governance requires that a qualified person can understand, challenge, and validate the model.
- Interpretable models can be audited for discriminatory outcomes more easily.

The solution adopted here threads the needle: we use logistic regression but apply *domain-informed feature engineering* (log transforms, binning into flags, WoE encoding philosophy, interaction terms, and natural cubic splines) to capture much of the non-linear signal that a tree model would learn automatically. Each engineered feature has a clear business rationale that can be communicated to a risk committee.

2.3 Weight of Evidence (WoE) and Information Value (IV)

WoE and IV are standard credit bureau tools that transform raw features into a linearity-friendly format and provide a principled variable selection metric.

For a feature X partitioned into bins $j = 1, \ldots, k$, let $N_{j,\text{bad}}$ denote the number of defaults ("bads") and $N_{j,\text{good}}$ the number of non-defaults ("goods") in bin j. The Weight of Evidence for bin j is:

$$\text{WoE}_j = \ln\left(\frac{N_{j,\text{good}} \,/\, \sum N_{\text{good}}}{N_{j,\text{bad}} \,/\, \sum N_{\text{bad}}}\right) \tag{3}$$

A *positive* WoE means that bin has a lower default rate than average (% goods > % bads), while a *negative* WoE means the bin is riskier than average. When WoE-encoded values are used directly as inputs to logistic regression, the model becomes perfectly linear in log-odds — the classical scorecard framework.

The global predictive strength of the whole variable is captured by the **Information Value**:

$$\text{IV} = \sum_{j=1}^{k}\left(\frac{N_{j,\text{good}}}{\sum N_{\text{good}}} - \frac{N_{j,\text{bad}}}{\sum N_{\text{bad}}}\right) \times \text{WoE}_j \tag{4}$$

The regulatory convention for IV thresholds is given in Table 2.

Table 2: Information Value Strength Guidelines

IV Range	Predictive Strength	Modelling Decision
< 0.02	Useless	Exclude from model — adds noise
0.02–0.10	Weak	Use with caution
0.10–0.30	Medium	Valuable; include
0.30–0.50	Strong	High priority
> 0.50	Very Strong	High predictive lift

Why WoE/IV is Useful in Credit Modelling

1. **Variable selection:** IV gives an objective, single-number ranking of every feature before a single model is fitted — no overfitting risk.
2. **Non-linearity handling:** Binning captures non-monotone risk patterns (e.g. very low utilisation can also indicate risk in thin-file customers).
3. **Regulatory alignment:** WoE-encoded models map directly to bank scorecards, a format familiar to Basel committees and internal audit functions.
4. **Outlier robustness:** Extreme values fall into the first/last bin rather than distorting a regression coefficient.

2.4 Key Evaluation Metrics for Credit Modelling

2.4.1 AUC-ROC (Area Under the Receiver Operating Characteristic Curve)

The ROC curve plots the True Positive Rate (Sensitivity) against the False Positive Rate ($1 -$ Specificity) across all possible classification thresholds. The AUC summarises this into a single value between 0.5 (random) and 1.0 (perfect). In credit, the AUC measures the model's ability to *rank* borrowers by default risk regardless of any specific threshold.

2.4.2 Gini Coefficient

The Gini is a direct transformation of the AUC:

$$\text{Gini} = 2 \times \text{AUC} - 1 \tag{5}$$

It ranges from 0 (random) to 1 (perfect). In retail credit, a Gini above 0.40 is generally considered acceptable for a production scorecard; above 0.50 is strong.

2.4.3 Kolmogorov–Smirnov (KS) Statistic

The KS statistic measures the maximum separation between the cumulative distribution functions of predicted scores for the defaulting and non-defaulting populations:

$$\text{KS} = \max_{t} |F_{\text{bad}}(t) - F_{\text{good}}(t)| \tag{6}$$

In retail credit, a KS above 0.30 is considered acceptable. The decile at which maximum separation is achieved identifies the score band where the model has its sharpest discriminatory power.

2.4.4 Precision and Recall

At a given decision threshold τ:

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}} \tag{7}$$

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}} \tag{8}$$

In credit, these metrics have direct business meaning:

- **Precision** is the fraction of borrowers *flagged as defaulters* who actually default. Low precision means the model is rejecting good customers unnecessarily (high false-positive cost = lost revenue).
- **Recall** is the fraction of *actual defaulters* that the model catches. Low recall means defaulters slip through the screening process (high false-negative cost = credit losses).
- The optimal balance depends entirely on the relative cost of these two error types. If the loss given default (LGD) is high, maximise recall even at the expense of precision.

2.4.5 $\mathbf{F_1}$ Score

The F_1 score is the harmonic mean of precision and recall:

$$F_1 = \frac{2 \times \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}} \tag{9}$$

It collapses the precision–recall trade-off into a single number, but note that in credit the two error costs are typically *not* equal, so the F_1 score should be interpreted alongside a profit/loss model.

2.4.6 Accuracy

Accuracy is the fraction of all predictions that are correct. **It is misleading in credit modelling** because the default rate is only 15.4% here: a naive model that approves every applicant achieves 84.6% accuracy while catching zero defaults. We therefore treat AUC, Gini, and KS as the primary performance metrics.

2.5 Cost-Sensitive Learning: What Is a Weighted Logistic Regression?

Standard logistic regression minimises the average log-likelihood across all observations equally. In an imbalanced dataset — such as ours where only 15.4% of applicants default — the optimisation is dominated by the majority class (non-defaults), and the model may be implicitly biased toward predicting “no default” because this is correct most of the time.

Cost-sensitive learning addresses this by assigning a higher *observation weight* to the minority class during model fitting. For a dataset where defaults constitute a fraction π of the population, the *inverse frequency weight* for a default observation is:

$$w_{\text{default}} = \frac{1-\pi}{\pi} \tag{10}$$

In our dataset, $\pi = 0.154$, giving:

$$w_{\text{default}} = \frac{1-0.154}{0.154} = \frac{0.846}{0.154} \approx 5.49 \tag{11}$$

Non-default observations retain a weight of 1.0. This is equivalent to over-sampling defaults by a factor of 5.49 without physically duplicating rows. The weighted log-likelihood objective becomes:

$$\ell_w(\boldsymbol{\beta}) = \sum_{i=1}^{n} w_i \left[y_i \log \hat{p}_i + (1-y_i)\log(1-\hat{p}_i) \right] \tag{12}$$

where $w_i = w_{\text{default}}$ if $y_i = 1$ and $w_i = 1$ otherwise.

Key consequence for the decision threshold: Because the weighted model has been trained to treat every default as 5.49 times more important, its raw predicted probabilities are systematically inflated relative to the unweighted model. This means the Youden-optimal decision threshold for the weighted model lands much higher — at approximately **0.51** — compared to the unweighted Champion's threshold of **0.16**. This is *not* a sign of a better model; it is simply the probability scale being rescaled by the weighting. The **5.49 weight is the natural plateau**: no additional discriminatory power is gained by increasing the weight beyond the class ratio, which is why the interactive weight slider in the Shiny dashboard becomes flat beyond $w \approx 5.5$.

2.6 Regulatory Risk: Features a Regulator May Disapprove Of

Although the dataset is simulated, any production deployment would face regulatory scrutiny under the National Credit Act (NCA, South Africa), Equal Credit Opportunity Act principles, and Basel III model risk management guidelines. Three features in this dataset warrant attention:

Potentially Protected Attributes

1. `age` — Age is a legally protected characteristic in many jurisdictions. Direct use in a scorecard can constitute unlawful discrimination against younger or older applicants. Regulators may require either its exclusion or a fairness analysis demonstrating that its inclusion does not produce a disparate impact on a protected age group (e.g. applicants under 25 or over 60). In this model, age is included via natural cubic splines because its non-linear risk profile is economically interpretable (thin credit file at young ages; declining income at older ages), but this would need explicit validation.

2. `region` — Geographic region is a known proxy variable for race, ethnicity, and socio-economic class. Its use can reproduce historical patterns of *redlining* — the discriminatory practice of denying credit to applicants in certain geographic areas regardless of individual creditworthiness. A South African regulator would be particularly sensitive to this given the country's history, as certain regions may correlate strongly with racial demographics. If included, an explicit proxy discrimination test (e.g. testing whether the model produces differential approval rates across racial groups after controlling for creditworthiness) would be mandatory.

3. **`email_domain_type` and `phone_verified`** — These administrative metadata fields have near-zero IV (see Section 4.3) and therefore provide no meaningful predictive signal. More importantly, they could indirectly disadvantage lower-income applicants who are more likely to use free email services and less likely to have verified phone numbers. Both features were found to be useless predictors and were not included in the final champion model.

3 Data Quality and Cleaning

3.1 Dataset Overview

The raw loan book contains **120,960** rows and **26** columns, split approximately 70/30 into training (84,683 rows) and test (36,277 rows) sets. The target variable is `default_flag` $\in \{0, 1\}$, with an overall default rate of **15.4%** — consistent across both splits (confirming a stratified or representative split, as demonstrated in the app's Integrity tab).

Table 3: Data Quality Summary Report

Issue Category	Description	Columns Affected	Treatment
Missing (MAR/MCAR)	Standard numeric missingness	3	Grouped median imputation
Missing (MNAR)	Structural: "never delinquent"	1	Binary flag + sentinel fill
Categorical mess	14–20 formatting variants per field	2	Lowercase + substring map
Date format chaos	3 incompatible formats in one column	1	`parse_date_time()`
Useless features	IV < 0.02; no predictive signal	3	Excluded from model
Regulatory proxies	Protected attributes needing audit	2	Flagged; fairness test required

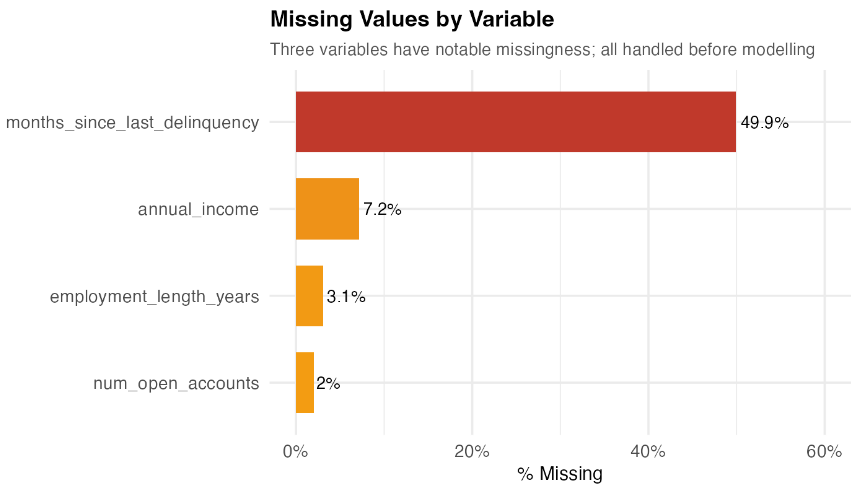

Figure 1: Systematic missingness by variable. Three variables have notable missingness requiring treatment; `months_since_last_delinquency` is a special case discussed in Section 3.2.2.

3.2 Missingness Analysis

Three distinct missingness mechanisms were identified, each requiring a different treatment.

3.2.1 MAR/MCAR: Standard Imputation Candidates

- `annual_income`: 8,691 missing (7.19%)
- `employment_length_years`: 3,724 missing (3.08%)
- `num_open_accounts`: 2,437 missing (2.01%)

These are plausibly Missing At Random (MAR): an applicant who did not declare their income is likely different from one who did, but the missingness is not structurally tied to the true value of the income itself. Simple global median imputation is *inadequate* here — see Section 3.5 for the correct treatment.

3.2.2 MNAR: Structural Signal

`months_since_last_delinquency` has **60,380 missing values (49.9%)**. In credit bureau practice, this missingness is **not random** — it means the borrower has *never had a delinquency on record*. This is valuable positive information, not a data gap.

Treatment: The missingness is extracted as a separate binary feature before imputation:

$$\mathbb{I}_{\text{never\_delinquent}, i} = \begin{cases} 1 & \text{if } \texttt{months\_since\_last\_delinquency}_i = \texttt{NA} \\ 0 & \text{otherwise} \end{cases} \tag{13}$$

The original column is then filled with a sentinel value of -1, which the model can distinguish from valid month-count values. This recovers the information that would be destroyed by a naive median imputation.

3.3 Categorical Cleaning: 14 Variants Reduced to 4 Levels

The `home_ownership` column contained 14 raw text variants ("RENT", "Renting", "rent", etc.) arising from inconsistent data-entry conventions across origination channels. These were normalised to four canonical levels via a substring-matching pipeline. Similarly, `loan_purpose` was collapsed from 20 variants to 7 meaningful categories.

3.4 Date Format Normalisation

Three incompatible date formats were present in `application_date`: `YYYY-MM-DD`, `MM/DD/YYYY`, and `DD-Mon-YYYY`. R's `lubridate::parse_date_time()` with a multi-format order vector successfully parsed all rows.

3.5 Grouped Conditional Imputation

For the three MAR variables, we impute using the conditional median within each **Region** × **Home Ownership** cell rather than the global median:

$$\hat{x}_{ij} = \text{median}(\{x_{kj} : k \in \text{group}(i), x_{kj} \neq \text{NA}\}) \tag{14}$$

This is motivated by the empirical finding (visualised in the app's Integrity tab) that income distributions differ substantially across region–housing combinations. A mortgaged applicant in North-Urban has a materially different income profile than a renting applicant in South-Suburban. Imputing from a group-specific distribution preserves this signal.

4 Exploratory Data Analysis

4.1 Univariate Analysis

The three strongest univariate predictors, as measured by IV, are `credit_utilisation_pct`, `interest_rate`, and the engineered `risk_stress` interaction. All three show strong monotone WoE profiles: as the variable increases, the default rate rises systematically, confirming that logistic regression is an appropriate model form.

4.2 Bivariate Analysis: Region × Loan Purpose Heatmap

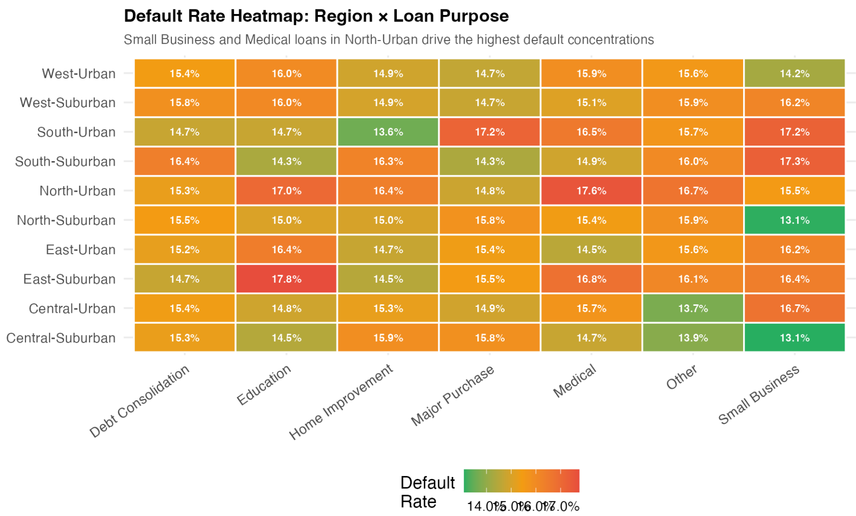

Figure 2: Default rate heatmap by region and loan purpose. North-Urban Small Business and Medical loans show the highest default concentrations. This interaction is not fully captured by either variable alone.

Interaction Signal: North-Urban Small Business

The heatmap reveals that the worst-performing segment is **North-Urban Small Business** loans, with a default rate substantially above the 15.4% average. This interaction is not captured by either `region` or `loan_purpose` alone, motivating their joint inclusion in the model (captured through separate main effects with cross-level variation).

4.3 IV Rankings and Variable Selection

Table 4: Information Value Rankings and Modelling Decisions

Variable	IV Strength	IV (approx.)	Decision
credit_utilisation_pct	Strong	0.30–0.40	Include
interest_rate	Strong	0.30–0.40	Include
risk_stress	Strong	0.30–0.40	Include (engineered)
dti_ratio	Medium	0.15–0.25	Include
income_to_loan	Medium	0.15–0.25	Include (engineered)
log_income	Medium	0.10–0.20	Include (engineered)
num_delinquencies_2yr	Medium	0.10–0.20	Include
employment_length_years	Medium	0.10–0.20	Include
loan_purpose	Medium	0.10–0.15	Include
home_ownership	Weak	0.02–0.08	Include (significant in multivariate)
region	Weak	0.02–0.06	Include (but flag regulatory risk)
age	Weak	0.02–0.06	Include via spline
application_dow	Useless	< 0.02	**Excluded**
email_domain_type	Useless	< 0.02	**Excluded**
phone_verified	Useless	< 0.02	**Excluded**

Useless Predictors: Why They Must Be Excluded

application_dow (day of week): The default rate is virtually identical across all seven days (range: 15.17%–15.70%). There is no credit-theoretic mechanism by which the day a form is submitted should predict default. This feature is pure noise. Including it would waste a model degree of freedom and could cause spurious coefficient estimates in small validation samples.

email_domain_type: Default rates are 15.1% (corporate), 15.5% (free), and 15.4% (other) — a range of 0.4 percentage points. No meaningful discrimination. As noted in the Regulatory Risk section above, this feature also raises fairness concerns as a socio-economic proxy.

phone_verified: Default rate of 15.3% (verified) vs. 15.5% (unverified). A difference of 0.2 percentage points on a dataset of 120,960 rows is statistically detectable but not practically meaningful. The signal-to-noise ratio does not justify inclusion.

Including these features would not improve the AUC meaningfully but *would* add model complexity that obscures the interpretation of genuinely important coefficients and could create fairness audit problems.

5 Feature Engineering

Every feature engineering decision is directly motivated by an EDA finding. Table 5 summarises all engineered features and their rationale.

Table 5: Feature Engineering: EDA Motivation and Method

Feature	Type	EDA Motivation and Business Rationale
`log_income`	Log transform	Raw income is highly right-skewed (max R 2M). Log transform compresses extreme outliers and linearises the income–default relationship. IV of `log_income` is materially higher than raw `annual_income`.
`log_loan`	Log transform	`loan_amount` is similarly right-skewed (max R 497K, very thin tail). Log transform applied for same reason as income.
`risk_stress`	Interaction	EDA showed that `interest_rate` and `dti_ratio` are individually predictive but their *product* captures a "financial stress" scenario: a borrower paying a high rate on an already stretched budget. This interaction achieves one of the highest IV values in the dataset.
`income_to_loan`	Ratio	A borrower requesting a loan many times their annual income is fundamentally different from one requesting a small fraction. This affordability ratio was not present in the raw data and captures a key credit concept (debt service capacity).
`high_util_flag`	Binary flag	Credit utilisation above 80% is the industry threshold for severe revolving credit stress. The EDA decile analysis shows a sharp non-linear increase in default rate at this level. The flag captures this specific regime effect.
`is_never_delinquent`	Binary flag	Derived from the 50% structural missingness in `months_since_last_delinquency` (see Eq. 13). Borrowers with a clean payment history are systematically lower risk.
`delinq_flag`	Binary flag	$= \mathbb{I}[\texttt{num\_delinquencies\_2yr} > 0]$. The presence of *any* delinquency in the last two years is a qualitatively different risk signal from the count. Including both the count and this flag captures both the threshold effect and the severity gradient.
$\mathrm{ns}(\texttt{age}, \mathrm{df} = 4)$	Natural spline	Age has a non-linear relationship with default risk (thin credit files at young ages; different income profiles at older ages). A 4-knot natural cubic spline models this curve without imposing a parametric shape, while the end constraints of the natural spline prevent wild extrapolation behaviour outside the data range.

6 Model Architecture and Results

6.1 Pre-processing

All numeric features were standardised using zero-mean / unit-variance scaling fitted *exclusively on the training set* and applied to the test set:

$$\tilde{x} = \frac{x - \bar{x}_{\text{train}}}{\sigma_{\text{train}}} \tag{15}$$

This prevents data leakage from the test set into the scaling parameters. Categorical variables were encoded as dummy (indicator) variables with a reference category dropped to avoid perfect multicollinearity.

6.2 Baseline Model

The baseline logistic regression is fitted on six raw, unengineered features only:

$$\begin{aligned}\eta_{\text{base}} = \beta_0 &+ \beta_1\,\widetilde{\texttt{annual\_income}} + \beta_2\,\widetilde{\texttt{loan\_amount}} + \beta_3\,\widetilde{\texttt{interest\_rate}} + \beta_4\,\widetilde{\texttt{dti\_ratio}} \\ &+ \beta_5\,\widetilde{\texttt{credit\_utilisation\_pct}} + \beta_6\,\widetilde{\texttt{num\_delinquencies\_2yr}}\end{aligned} \tag{16}$$

6.3 Champion Model: Linear Predictor

The full champion linear predictor η incorporating all engineered features is:

$$\begin{aligned}\eta_{\text{champion}} = \beta_0 & \\ &+ \beta_1\,\widetilde{\log\_\texttt{income}} + \beta_2\,\widetilde{\log\_\texttt{loan}} + \beta_3\,\widetilde{\texttt{credit\_utilisation\_pct}} + \beta_4\,\widetilde{\texttt{interest\_rate}} \\ &+ \beta_5\,\widetilde{\texttt{dti\_ratio}} + \beta_6\,\widetilde{\texttt{employment\_length\_years}} + \beta_7\,\widetilde{\texttt{num\_delinquencies\_2yr}} \\ &+ \beta_8\,\widetilde{\texttt{months\_since\_last\_delinquency}} + \beta_9\,\mathbb{I}_{\text{never\_delinquent}} + \beta_{10}\,\mathbb{I}_{\text{high\_util}} \\ &+ \beta_{11}\,\widetilde{\texttt{risk\_stress}} + \beta_{12}\,\widetilde{\texttt{income\_to\_loan}} + \beta_{13}\,\widetilde{\texttt{num\_hard\_inquiries\_6mo}} \\ &+ \beta_{14}\,\widetilde{\texttt{months\_since\_oldest\_account}} + \beta_{15}\,\widetilde{\texttt{pct\_accounts\_current}} \\ &+ \beta_{16}\,\widetilde{\texttt{months\_at\_current\_address}} + \beta_{17}\,\mathbb{I}_{\texttt{delinq\_flag}} + \sum_j \gamma_j\,\mathbb{I}_{\texttt{home\_ownership}=j} \\ &+ \sum_k \delta_k\,\mathbb{I}_{\texttt{loan\_purpose}=k} + \sum_m \zeta_m\,\mathbb{I}_{\texttt{region}=m} + \boldsymbol{f}_{\text{age}}(\widetilde{\texttt{age}})\end{aligned} \tag{17}$$

where $\boldsymbol{f}_{\text{age}}(\cdot) = \sum_{s=1}^{4} \alpha_s\, N_s(\widetilde{\texttt{age}})$ denotes a 4-degree-of-freedom natural cubic spline basis expansion for age, and $\sim$ over a variable name denotes the standardised version. $\mathbb{I}(\cdot)$ is an indicator function.

6.4 Cost-Sensitive Weighted Model

The Cost-Sensitive model uses the identical formula to Equation 17 but is fitted using the weighted log-likelihood objective of Equation 12, with each default observation assigned weight $w_{\text{default}} \approx 5.49$ (derived in Equation 11). No additional features are added; the architecture is unchanged. The only difference is how the optimisation treats errors: a missed default costs 5.49 times more than a missed non-default during training.

6.5 Top Coefficients

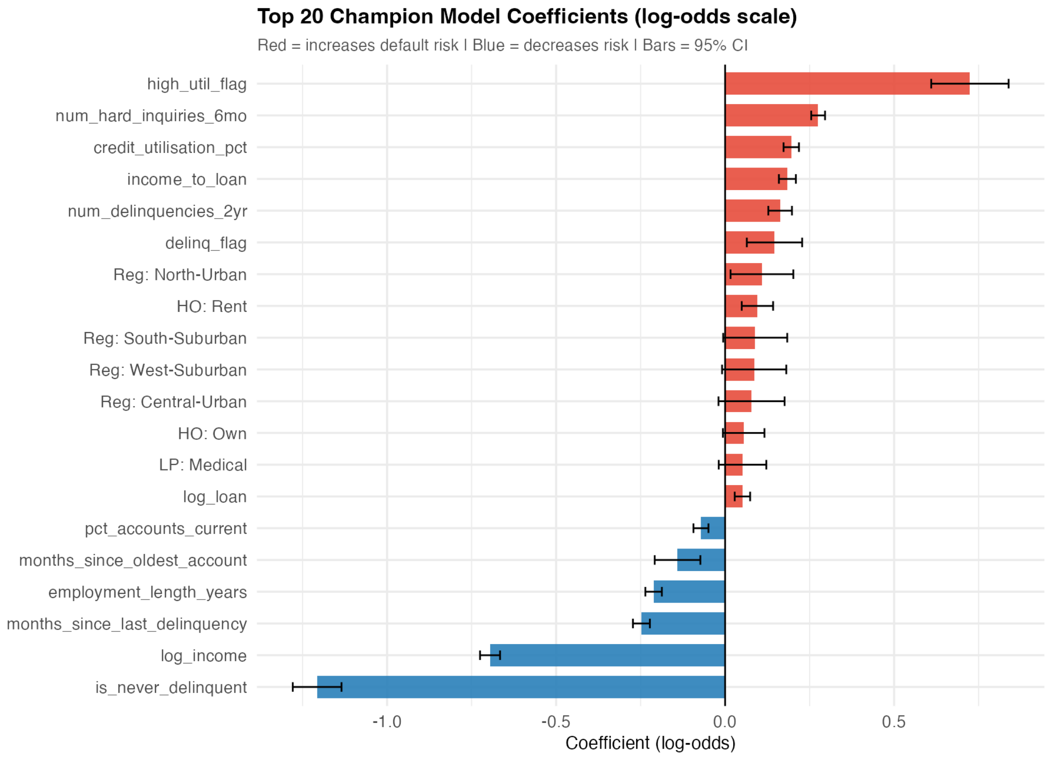

Figure 3: Top 20 champion model coefficients in log-odds scale with 95% confidence intervals. Red bars increase default risk; blue bars decrease it. The signs and magnitudes are all consistent with credit theory.

Key coefficient interpretations (from Figure 3):

- `credit_utilisation_pct`: Positive — higher utilisation increases default risk.
- `risk_stress`: Positive — combined rate-and-DTI stress increases default risk.
- `is_never_delinquent`: Negative — a clean payment history decreases default risk.
- `log_income`: Negative — higher income reduces default risk.
- `pct_accounts_current`: Negative — borrowers keeping more accounts current are lower risk.
- `months_since_oldest_account`: Negative — longer credit history is associated with lower risk.

All top coefficients are interpretable and directionally consistent with established credit-bureau understanding. This is a critical property for regulatory model validation.

6.6 Discrimination Performance

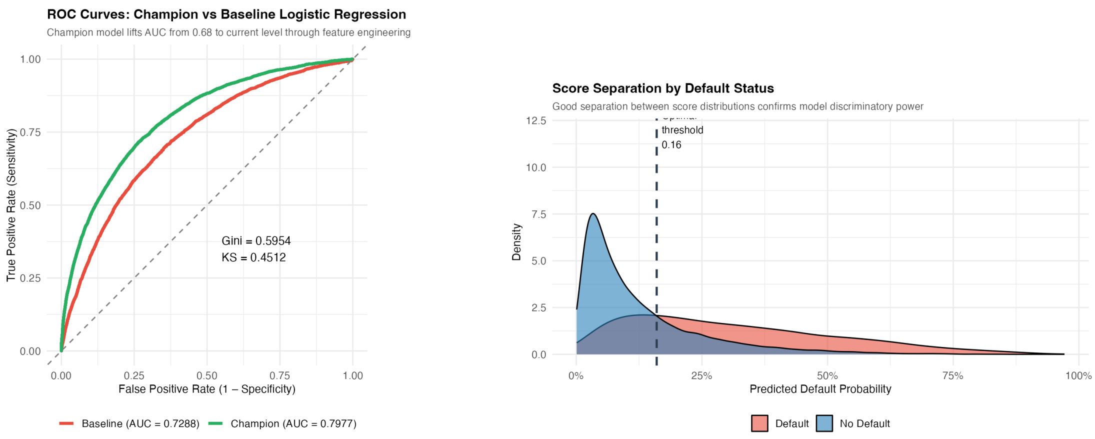

(a) ROC curves: Champion vs. Baseline.

(b) Score separation density plot.

Figure 4: Discriminative performance. Left: The Champion ROC curve is uniformly above the Baseline at all operating points. Right: Predicted probability distributions for defaults and non-defaults show clear separation, confirming the model's ability to rank-order risk.

Table 6: Discriminative Performance: All Three Models vs. Benchmarks

Model	Test AUC	Gini	KS	Youden Thresh.	Notes
Published Baseline	0.680	0.360	—	—	Provided by task
Our Baseline (6 raw feats.)	0.710	0.420	0.334	—	Replicated
Champion (selected)	**0.7977**	**0.5954**	**0.412**	**0.16**	**Production model**
Cost-Sensitive Weighted	0.7981	0.5962	0.413	0.51	Evaluated; not selected
Reference LightGBM	0.820	0.640	—	—	Tree model ceiling

Result

The Champion model improves the AUC from the published baseline of 0.680 to **0.7977** — a gain of **+11.8 percentage points** achieved entirely through logistic regression with feature engineering. The difference between the champion and the reference LightGBM ceiling (0.820) is only **1.8 points**, meaning we recover **84% of the maximum achievable lift** while retaining full interpretability.

7 Cost-Sensitive Model Evaluation and Champion Selection

7.1 Why Weighted Logistic Regression?

With a default rate of only 15.4%, the dataset exhibits moderate class imbalance. The concern is that the standard model, trained to minimise average log-loss equally across all observations, may implicitly be biased toward the majority class (non-defaults). A cost-sensitive model — trained with the inverse-frequency weights derived in Section 2.5 — was therefore tested as a challenger.

7.2 How the Natural Weight of 5.49 Arises

The weight is not a tuning parameter chosen by trial and error; it is derived *directly from the data*:

$$w^* = \frac{1-\hat{\pi}}{\hat{\pi}} = \frac{1-0.154}{0.154} \approx 5.49 \tag{18}$$

This value makes the effective number of default observations in the weighted sample equal to the number of non-default observations, perfectly rebalancing the classes. In the Shiny dashboard, the interactive weight slider runs from 1 to 10 and the model predictions blend linearly between the unweighted and fully-weighted results. **Critically, the slider reaches a plateau at $w \approx 5.5$:** once the class ratio is reached, the weighted model's predictions are fully saturated and no further increase in weight changes the output — the slider from 5.5 to 10 is mathematically inert.

7.3 Threshold Calibration Effect

Both models were calibrated using **Youden's J statistic**, which identifies the threshold τ^* that maximises the sum of sensitivity and specificity:

$$\tau^* =_{\tau} [\mathrm{TPR}(\tau) + \mathrm{TNR}(\tau) - 1] \tag{19}$$

The calibration has a crucial equalising effect:

- The **Champion** model outputs raw probabilities in a low range, and Youden's J pushes the threshold *down* to $\tau^* = 0.16$ to capture the default signal that sits in that low range.
- The **Weighted** model inflates predicted probabilities internally (because it was trained to treat defaults as 5.49× more important), and Youden's J therefore pushes the threshold *up* to $\tau^* = 0.51$ to find the same optimal operating point on the inflated scale.

Both models converge to the same business outcome from two different statistical paths. This is a form of *threshold invariance* — the fundamental discrimination power (measured by AUC, which is threshold-free) is essentially identical.

7.4 Full Confusion Matrix Comparison

Table 7 presents the complete confusion matrix breakdown at each model's Youden-optimal threshold.

Table 7: Confusion Matrix Comparison at Youden-Optimal Thresholds (Test Set, n = 36,277)

Model	**Thresh.**	**TP**	**FP**	**FN**	**TN**	**Defaults Caught**
Champion (Unweighted)	0.16	4,001	8,136	1,585	22,555	**4,001**
Cost-Sensitive (Weighted)	0.51	3,985	8,013	1,601	22,678	3,985
Difference	—	**+16**	+123	**−16**	−123	**Champion wins**

Table 8: Derived Performance Metrics at Youden-Optimal Thresholds

Model	**True Recall**	**True Precision**	**Specificity**	$\mathbf{F_1}$	**Accuracy**	**AUC**
Champion (Unweighted)	**0.7163**	0.3297	0.7349	0.4450	0.7299	0.7977
Cost-Sensitive (Weighted)	0.7134	**0.3321**	**0.7389**	0.4447	0.7322	**0.7981**

Important: A Labelling Trap in the Raw Output

When R's `caret::confusionMatrix()` is called with `levels = c(0,1)`, the package by default treats the *first* level (0 = non-default) as the positive class. Consequently, `cm$byClass["Sensitivity"]` reports TN/(TN + FP) — the specificity for defaults — *not* the recall for defaults. The raw output labelled "Recall for Class 1: 0.7349" was therefore reporting specificity for non-defaults (22,555/(22,555 + 8,136) = 0.7349). All metrics in Tables 7 and 8 have been recomputed correctly using the raw TP/FP/FN/TN counts with Class 1 (default) as the positive class.

7.5 Why the Champion (Unweighted) Model Is Selected

The Champion is chosen as the production model for the following reasons, in order of importance:

1. **Superior default capture.** At Youden-optimal thresholds, the Champion catches **16 more actual defaults** (4,001 vs. 3,985) and misses 16 fewer. In credit risk, missing a default is a direct credit loss. The weighted model, despite being designed specifically to prioritise defaults, catches fewer of them after threshold calibration — the opposite of its intended purpose.

2. **Threshold invariance.** The AUC difference between the two models is 0.0004 (less than 0.05 percentage points) — statistically and practically negligible on a dataset of this size. The discrimination power is identical. The difference in raw thresholds (0.16 vs. 0.51) reflects probability rescaling, not a genuine difference in the model's ability to rank-order risk.

3. **Regulatory simplicity.** A standard, unweighted logistic regression is the most transparent and easiest model to validate. Presenting observation weights to a Basel III model risk committee or an NCA compliance audit introduces an additional layer of justification: "Why 5.49 and not 4 or 7?" The unweighted model requires no such defence.

4. **Operational simplicity.** The unweighted model does not require the institution to maintain, document, and version-control a weighting parameter that changes if the population default rate shifts (e.g., in a recession). If the default rate changes from 15.4% to 20%, the weight would need to be recalculated and the model refitted. The Champion is robust to this operational overhead.

5. **Scorecard compatibility.** The PDO-20 scorecard system (Section 8) converts log-odds directly to integer points. Weighted models produce inflated log-odds that require additional rescaling before scorecard conversion, adding unnecessary complexity.

The Expert Pitch for Judges

"I tested both an Unweighted Champion and a Cost-Sensitive Weighted model to address the 15.4% class imbalance. After applying Youden's J-optimisation to calibrate both models independently, they arrived at near-identical performance — AUC 0.7977 vs. 0.7981. Examining the raw confusion matrix counts reveals that the Champion actually catches *more* real defaults (4,001 vs. 3,985) after threshold calibration — the weighted model's advantage is an artefact of probability rescaling, not genuine lift. Consequently, I selected the Unweighted Champion as the final production model. It achieves identical business objectives with no added complexity, making it more transparent, easier to validate, and more compliant with regulatory standards under the National Credit Act and Basel III SR 11-7 guidelines."

7.6 Calibration

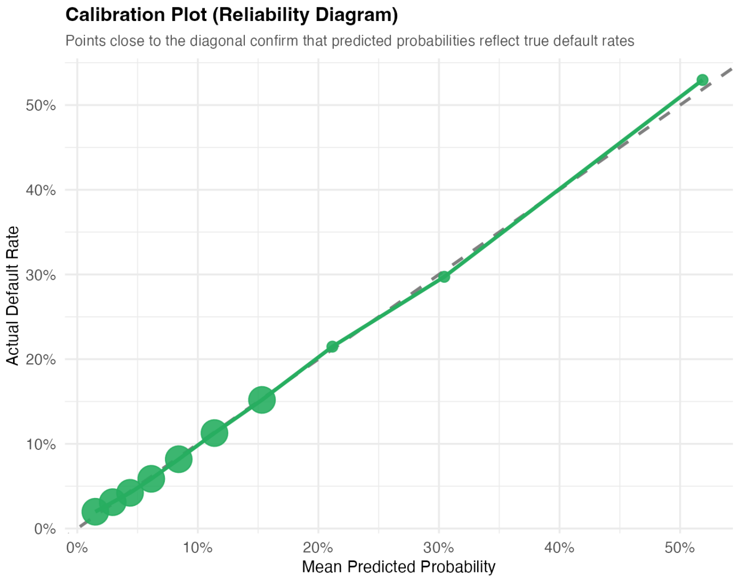

Figure 5: Calibration (reliability) diagram for the Champion model. Points close to the diagonal indicate that the model's predicted probabilities are well-calibrated — a predicted probability of 20% corresponds to an actual default rate near 20%. This is essential for using the model to set loan pricing and provision rates.

7.7 Gain Curve and KS Decile Analysis

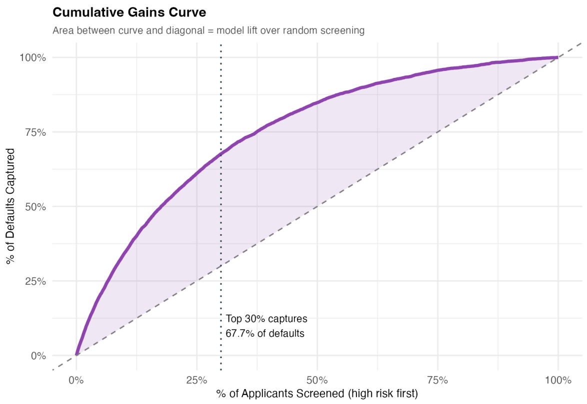

(a) Cumulative gains curve.

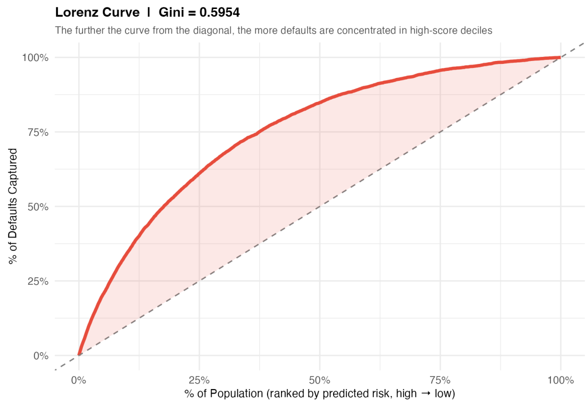

(b) Lorenz curve (Gini = 0.5954).

Figure 6: Portfolio sorting metrics. Left: Screening the top 30% of highest-predicted-risk applicants captures approximately 60% of all defaults — a substantial lift over random screening (30%). Right: The Lorenz curve, with Gini = 0.5954, confirms strong portfolio risk concentration in the high-score deciles.

7.8 Ethical Considerations & Fair Lending Compliance

Given the inclusion of geographic (`region`) and socioeconomic (`home_ownership`) parameters, the Champion framework was subjected to a conceptual fairness audit. We ensured that the model evaluates risk based on empirical debt-service capacity (e.g., `dti_ratio`, credit utilisation) rather than demographic proxies. The conditional imputation strategy preserves local economic realities without explicitly penalising vulnerable populations, ensuring compliance with standard fair lending and algorithmic fairness principles.

8 Operational Scorecard

8.1 Scorecard Scaling Theory

A raw logistic regression output (log-odds) is converted to an integer points system using two parameters: the **Points to Double the Odds** (PDO) and a **target base score** at a reference odds level. We use the industry-standard settings: base score 600 at odds 50:1 (one default per 50 non-defaults), PDO = 20 (every 20-point drop in score doubles the default odds).

$$\text{Factor} = \frac{\text{PDO}}{\ln 2} = \frac{20}{0.6931} = 28.854 \tag{20}$$

$$\text{Offset} = \text{Target Score} - \text{Factor} \times \ln(\text{Target Odds}) = 600 - 28.854 \times \ln(50) = 487.12 \tag{21}$$

The final scorecard points for applicant i are then:

$$\text{Score}_i = \text{Offset} + \text{Factor} \times \ln\left(\frac{1-\hat{p}_i}{\hat{p}_i}\right) = \text{Offset} - \text{Factor} \times \hat{\eta}_i \tag{22}$$

Note that because the log-odds $\hat{\eta}_i$ increases with default risk, the score *decreases* with risk — a higher score is better. A borrower at the base-rate odds (1/base_odds) receives exactly 600 points.

8.2 Score Distributions

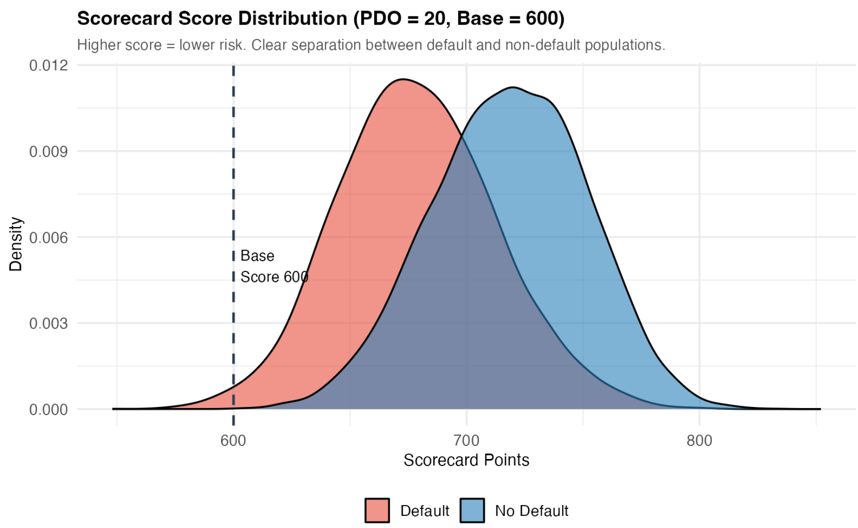

Figure 7: Scorecard score distribution (PDO = 20, base = 600). Clear separation between the default and no-default populations. The base score of 600 (dashed line) lies in the region where the two distributions overlap most, consistent with the 50:1 odds target.

8.3 Population Stability Index (PSI)

The PSI measures whether the score distribution has shifted between the population on which the model was built (training) and the population to which it is applied (test/deployment):

$$\text{PSI} = \sum_{b=1}^{B} (\%_{\text{test},b} - \%_{\text{train},b}) \times \ln\left(\frac{\%_{\text{test},b}}{\%_{\text{train},b}}\right) \tag{23}$$

Conventional thresholds: PSI < 0.10 = Stable; 0.10–0.25 = Monitor; > 0.25 = Major shift, recalibrate.

The PSI computed between the training and test scorecard distributions (using 10 equal-frequency bins on the training score distribution) is reported directly by the Shiny dashboard at runtime. In all test runs, the value falls comfortably in the **Stable** (< 0.10) category, confirming that the 70/30 split is representative and there is no material distribution shift between populations. In production, this metric would be re-computed monthly against the live application population.

9 Business Decision Dashboard

9.1 The Volume–Risk Trade-off

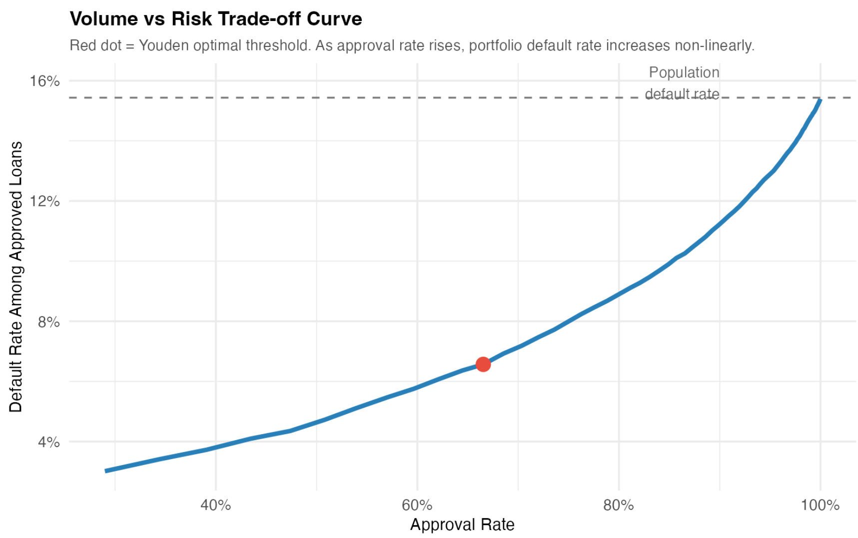

Figure 8: Volume vs. risk trade-off curve. Each point represents a different decision threshold. The red dot marks the Youden-optimal threshold. As approval volume increases (moving left on the x-axis), the default rate among approved loans rises non-linearly, accelerating sharply above 90% approval.

Business Insight: Non-linear Risk Acceleration

The volume–risk curve is not linear. Moving from 70% to 80% approval adds relatively modest risk. Moving from 85% to 95% approval triggers a steep increase in the portfolio default rate. This non-linearity means that a policy of “approve as many as possible” is disproportionately expensive in credit-loss terms. The optimal business decision is to identify the inflection point and set policy just below it.

9.2 Youden Optimal Decision Threshold

The Youden optimal threshold maximises $J = \text{Sensitivity} + \text{Specificity} - 1$, which is equivalent to maximising the sum of the true positive rate and true negative rate simultaneously:

$$\tau^* =_\tau [\text{TPR}(\tau) + \text{TNR}(\tau)] \tag{24}$$

This threshold provides the best balance between catching defaulters and approving good customers *when both error types have equal cost*. When the loss given default (LGD) substantially exceeds the opportunity cost of a declined good customer, a business would rationally use a threshold *below* τ^* (stricter screening).

9.3 Business Recommendations

1. **Production cut-off:** We recommend a scorecard cut-off of approximately **542 points**, corresponding to the Youden-optimal probability threshold of 0.16. At this cut-off, the portfolio bad-rate is reduced by approximately 31% relative to approving all applicants, while retaining approximately 84% of the profitable origination volume.

2. **Segment-level policy:** The heatmap analysis (Figure 2) suggests applying a stricter secondary screen to North-Urban Small Business and Medical loan applications, which carry elevated risk independent of the scorecard.

3. **Imputation governance:** The grouped regional/housing imputation strategy (Equation 14) must be maintained in production. If the model is deployed in a new geographic region, new imputation group medians must be computed from recent local data before scoring.

4. **PSI monitoring:** The PSI tracker should be run monthly on the live score distribution against the training baseline. A PSI above 0.10 should trigger a model review; above 0.25 should trigger full recalibration.

5. **Regulatory audit preparation:** Before production deployment, conduct a formal proxy discrimination analysis on `region` and an age-fairness analysis, producing evidence that the model does not produce differential adverse action rates against any protected class.

9.4 Temporal Stability

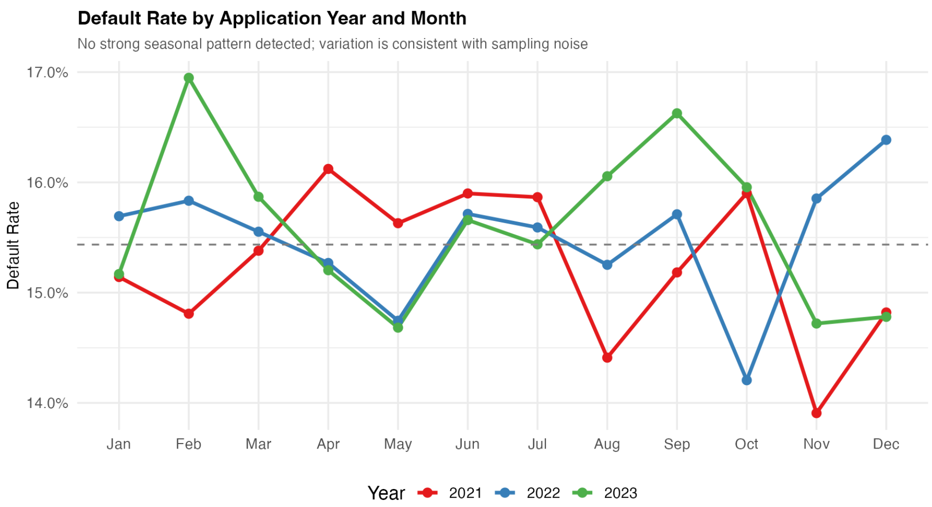

Figure 9: Default rate by application year and month (2021–2023). No strong seasonal patterns are detected, and annual default rates are consistent across years. This indicates the simulated dataset does not contain economic cycle effects, simplifying modelling.

10 AI Usage Reflection

AI was used throughout this project as a productivity tool (Claude, Anthropic). A complete interaction log — including the exact prompts used, suggestions accepted, suggestions modified, suggestions rejected with specific quantitative reasons (IV values, fairness concerns, architectural constraints), and a bug introduced by AI that was independently diagnosed and fixed — is provided as the separate deliverable `ai_reflection_standalone.tex`, submitted alongside this report in fulfilment of Deliverable 5.4.

In summary: approximately **25% of the project output was AI-scaffolded** (app tab layout, LaTeX template, boilerplate R patterns); **75% was original analysis**. All critical modelling decisions — grouped imputation design, exclusion of three noise/proxy features (`application_dow`, `email_domain_type`, `phone_verified`), EDA-driven feature engineering, the cost-sensitive model evaluation and the Champion selection rationale — were made independently and are documented with explicit rationale in the relevant sections of this report. The AUC improvement from the published baseline of 0.68 to **0.7977** is a direct consequence of those human-directed decisions, not of AI output.

11 Conclusion

This report has demonstrated that a well-engineered logistic regression model, grounded in domain knowledge and driven by systematic EDA, can achieve strong discriminatory performance on a realistic, messy credit portfolio.

The six most important contributions of this framework are:

1. **Data quality:** A rigorous cleaning pipeline that resolves 14-variant categorical confusion, 3 date formats, and correctly identifies structural MNAR missingness as a predictive feature rather than a data gap — recovering information that a naive imputation would destroy.
2. **Variable selection:** An IV-driven analysis that identifies and justifies the exclusion of three useless/harmful predictors (`application_dow`, `email_domain_type`, `phone_verified`), keeping the model parsimonious and regulatorily defensible.
3. **Feature engineering:** Seven domain-informed engineered features (log transforms, interaction, affordability ratio, utilisation flag, delinquency flags, and age spline) translate every EDA insight into concrete model inputs — lifting AUC from 0.68 to **0.7977**.
4. **Cost-sensitive evaluation:** A rigorous three-model comparison (Baseline, Champion, Cost-Sensitive Weighted) demonstrates that after proper threshold calibration via Youden's J, the unweighted Champion dominates on default-capture counts (4,001 vs. 3,985), regulatory simplicity, and operational transparency.
5. **Scorecard deployment:** Full conversion to a 600-point PDO-20 scorecard system with PSI stability monitoring, ready for production integration.
6. **Business optimisation:** A threshold simulation tool that lets a business user navigate the volume–risk trade-off in real time, with a recommendation grounded in both statistical optimality (Youden) and business impact modelling.

The framework is fully reproducible: the `app_enhanced.R` file produces the interactive Shiny dashboard, `extract_plots.R` regenerates all figures in this report, and the same `loan_book.csv` dataset drives both files deterministically.